

# Chapter Outline

---

Relational Model Concepts

Relational Model Constraints and Relational Database Schemas

Update Operations and Dealing with Constraint Violations

# Relational Model Concepts

---

The relational Model of Data is based on the concept of a *Relation*

- The strength of the relational approach to data management comes from the formal foundation provided by the theory of relations

We review the essentials of the *formal relational model* in this chapter

In *practice*, there is a *standard model* based on SQL – this is described in Chapters 8 and 9

Note: There are several important differences between the *formal* model and the *practical* model, as we shall see

# Relational Model Concepts

---

A Relation is a mathematical concept based on the ideas of sets

The model was first proposed by Dr. E.F. Codd of IBM Research in 1970 in the following paper:

- "A Relational Model for Large Shared Data Banks,"  
Communications of the ACM, June 1970

The above paper caused a major revolution in the field of database management and earned Dr. Codd the coveted ACM Turing Award

# Informal Definitions

---

Informally, a **relation** looks like a **table** of values.

A relation typically contains a **set of rows**.

The data elements in each **row** represent certain facts that correspond to a real-world **entity** or **relationship**

- In the formal model, rows are called **tuples**

Each **column** has a column header that gives an indication of the meaning of the data items in that column

- In the formal model, the column header is called an **attribute name** (or just **attribute**)

# Example of a Relation

---

# Informal Definitions

---

## Key of a Relation:

- Each row has a value of a data item (or set of items) that uniquely identifies that row in the table
  - Called the *key*
- In the STUDENT table, SSN is the key
- Sometimes row-ids or sequential numbers are assigned as keys to identify the rows in a table
  - Called *artificial key* or *surrogate key*

# Formal Definitions - Schema

---

The **Schema** (or description) of a Relation:

- Denoted by  $R(A_1, A_2, \dots, A_n)$
- $R$  is the **name** of the relation
- The **attributes** of the relation are  $A_1, A_2, \dots, A_n$

Example:

CUSTOMER (Cust-id, Cust-name, Address, Phone#)

- CUSTOMER is the relation name
- Defined over the four attributes: Cust-id, Cust-name, Address, Phone#

Each attribute has a **domain** or a set of valid values.

- For example, the domain of Cust-id is 6 digit numbers.

# Formal Definitions - Tuple

---

A **tuple** is an ordered set of values (enclosed in angled brackets '< ... >')

Each value is derived from an appropriate *domain*.

A row in the CUSTOMER relation is a 4-tuple and would consist of four values, for example:

- <632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">
- This is called a 4-tuple as it has 4 values
- A tuple (row) in the CUSTOMER relation.

A relation is a **set** of such tuples (rows)



# Formal Definitions - Domain

A **domain** has a logical definition:

---

- Example: “USA\_phone\_numbers” are the set of 10 digit phone numbers valid in the U.S.

A domain also has a data-type or a format defined for it.

- The USA\_phone\_numbers may have a format: (ddd)ddd-dddd where each d is a decimal digit.
- Dates have various formats such as year, month, date formatted as yyyy-mm-dd, or as dd mm,yyyy etc.

The attribute name designates the role played by a domain in a relation:

- Used to interpret the meaning of the data elements corresponding to that attribute
- Example: The domain Date may be used to define two attributes named “Invoice-date” and “Payment-date” with different meanings

# Formal Definitions - State

---

The **relation state** is a subset of the Cartesian product of the domains of its attributes

- each domain contains the set of all possible values the attribute can take.

Example: attribute Cust-name is defined over the domain of character strings of maximum length 25

- $\text{dom}(\text{Cust-name})$  is `varchar(25)`

The role these strings play in the CUSTOMER relation is that of the *name of a customer*.

# Formal Definitions - Summary

---

Formally,

- Given  $R(A_1, A_2, \dots, A_n)$
- $r(R) \subset \text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n)$

$R(A_1, A_2, \dots, A_n)$  is the **schema** of the relation

$R$  is the **name** of the relation

$A_1, A_2, \dots, A_n$  are the **attributes** of the relation

$r(R)$ : a specific **state** (or "value" or "population") of relation  $R$  – this is a *set of tuples* (rows)

- $r(R) = \{t_1, t_2, \dots, t_n\}$  where each  $t_i$  is an  $n$ -tuple
- $t_i = \langle v_1, v_2, \dots, v_n \rangle$  where each  $v_j$  *element-of*  $\text{dom}(A_j)$

# Formal Definitions - Example

---

Let  $R(A1, A2)$  be a relation schema:

- Let  $\text{dom}(A1) = \{0,1\}$
- Let  $\text{dom}(A2) = \{a,b,c\}$

Then:  $\text{dom}(A1) \times \text{dom}(A2)$  is all possible combinations:

$\{ \langle 0,a \rangle , \langle 0,b \rangle , \langle 0,c \rangle , \langle 1,a \rangle , \langle 1,b \rangle , \langle 1,c \rangle \}$

The relation state  $r(R) \subseteq \text{dom}(A1) \times \text{dom}(A2)$

For example:  $r(R)$  could be  $\{ \langle 0,a \rangle , \langle 0,b \rangle , \langle 1,c \rangle \}$

- this is one possible state (or “population” or “extension”)  $r$  of the relation  $R$ , defined over  $A1$  and  $A2$ .
- It has three 2-tuples:  $\langle 0,a \rangle , \langle 0,b \rangle , \langle 1,c \rangle$

# Definition Summary

<u>Informal Terms</u>		<u>Formal Terms</u>
Table		Relation
Column Header		Attribute
All possible Column Values		Domain
Row		Tuple
Table Definition		Schema of a Relation
Populated Table		State of the Relation

# Example – A relation

## STUDENT

---

# Characteristics Of Relations

---

Ordering of tuples in a relation  $r(R)$ :

- The tuples are *not considered to be ordered*, even though they appear to be in the tabular form.

Ordering of attributes in a relation schema  $R$  (and of values within each tuple):

- We will consider the attributes in  $R(A_1, A_2, \dots, A_n)$  and the values in  $t = \langle v_1, v_2, \dots, v_n \rangle$  to be ordered .
- (However, a more general alternative definition of relation does not require this ordering).

same state as previous figure  
(but with different order of  
tuples)

---



# Characteristics Of Relations

---

Values in a tuple:

- All values are considered atomic (indivisible).
- Each value in a tuple must be from the domain of the attribute for that column
  - If tuple  $t = \langle v_1, v_2, \dots, v_n \rangle$  is a tuple (row) in the relation state  $r$  of  $R(A_1, A_2, \dots, A_n)$
  - Then each  $v_i$  must be a value from  $dom(A_i)$
- A special **null** value is used to represent values that are unknown or inapplicable to certain tuples.

# Characteristics Of Relations

---

## Notation:

- We refer to **component values** of a tuple  $t$  by:
  - $t[A_i]$  or  $t.A_i$
  - This is the value  $v_i$  of attribute  $A_i$  for tuple  $t$
- Similarly,  $t[A_u, A_v, \dots, A_w]$  refers to the subtuple of  $t$  containing the values of attributes  $A_u, A_v, \dots, A_w$ , respectively in  $t$

# Relational Integrity Constraints

---

Constraints are **conditions** that must hold on **all** valid relation states.

There are three *main types* of constraints in the relational model:

- **Key** constraints
- **Entity integrity** constraints
- **Referential integrity** constraints

Another implicit constraint is the **domain** constraint

- Every value in a tuple must be from the *domain of its attribute* (or it could be **null**, if allowed for that attribute)

# Key Constraints

---

## Superkey of R:

- Is a set of attributes SK of R with the following condition:
  - No two tuples in any valid relation state  $r(R)$  will have the same value for SK
  - That is, for any distinct tuples  $t_1$  and  $t_2$  in  $r(R)$ ,  $t_1[SK] \neq t_2[SK]$
  - This condition must hold in *any valid state*  $r(R)$

## Key of R:

- A "minimal" superkey
- That is, a key is a superkey K such that removal of any attribute from K results in a set of attributes that is not a superkey (does not possess the superkey uniqueness property)

# Key Constraints (continued)

---

Example: Consider the CAR relation schema:

- CAR(State, Reg#, SerialNo, Make, Model, Year)
- CAR has two keys:
  - Key1 = {State, Reg#}
  - Key2 = {SerialNo}
- Both are also superkeys of CAR
- {SerialNo, Make} is a superkey but *not* a key.

In general:

- Any *key* is a *superkey* (but not vice versa)
- Any set of attributes that *includes a key* is a *superkey*
- A *minimal* superkey is also a key

# Key Constraints (continued)

---

If a relation has several **candidate keys**, one is chosen arbitrarily to be the **primary key**.

- The primary key attributes are underlined.

Example: Consider the CAR relation schema:

- CAR(State, Reg#, SerialNo, Make, Model, Year)
- We chose SerialNo as the primary key

The primary key value is used to *uniquely identify* each tuple in a relation

- Provides the tuple identity

Also used to *reference* the tuple from another tuple

- General rule: Choose as primary key the smallest of the candidate keys (in terms of size)
- Not always applicable – choice is sometimes subjective

# CAR table with two candidate keys – LicenseNumber chosen as Primary Key

---

# Relational Database Schema

---

## Relational Database Schema:

- A set  $S$  of relation schemas that belong to the same database.
- $S$  is the name of the whole **database schema**
- $S = \{R_1, R_2, \dots, R_n\}$
- $R_1, R_2, \dots, R_n$  are the names of the individual **relation schemas** within the database  $S$

Following slide shows a COMPANY database schema with 6 relation schemas



# COMPANY Database Schema

# Entity Integrity

---

## Entity Integrity:

- The *primary key attributes* PK of each relation schema R in S cannot have null values in any tuple of  $r(R)$ .
  - This is because primary key values are used to *identify* the individual tuples.
  - $t[PK] \neq \text{null}$  for any tuple  $t$  in  $r(R)$
  - If PK has several attributes, null is not allowed in any of these attributes
- Note: Other attributes of R may be constrained to disallow null values, even though they are not members of the primary key.

# Referential Integrity

---

A constraint involving **two** relations

- The previous constraints involve a single relation.

Used to specify a **relationship** among tuples in two relations:

- The **referencing relation** and the **referenced relation**.

# Referential Integrity

---

Tuples in the **referencing relation** R1 have attributes FK (called **foreign key** attributes) that reference the primary key attributes PK of the **referenced relation** R2.

- A tuple t1 in R1 is said to **reference** a tuple t2 in R2 if  $t1[FK] = t2[PK]$ .

A referential integrity constraint can be displayed in a relational database schema as a directed arc from R1.FK to R2.

# Referential Integrity (or foreign key) Constraint

---

## Statement of the constraint

- The value in the foreign key column (or columns) FK of the the **referencing relation** R1 can be **either**:
  - (1) a value of an existing primary key value of a corresponding primary key PK in the **referenced relation** R2, or
  - (2) a **null**.

In case (2), the FK in R1 should **not** be a part of its own primary key.

# Displaying a relational database schema and its constraints

---

Each relation schema can be displayed as a row of attribute names

The name of the relation is written above the attribute names

The primary key attribute (or attributes) will be underlined

A foreign key (referential integrity) constraints is displayed as a directed arc (arrow) from the foreign key attributes to the referenced table

- Can also point the the primary key of the referenced relation for clarity

Next slide shows the COMPANY **relational schema diagram**

# Referential Integrity Constraints for COMPANY database

# Other Types of Constraints

---

## Semantic Integrity Constraints:

- based on application semantics and cannot be expressed by the model per se
- Example: “the max. no. of hours per employee for all projects he or she works on is 56 hrs per week”

A **constraint specification** language may have to be used to express these

SQL-99 allows triggers and **ASSERTIONS** to express for some of these



# Populated database state

---

Each *relation* will have many tuples in its current relation state

The *relational database state* is a union of all the individual relation states

Whenever the database is changed, a new state arises

Basic operations for changing the database:

- INSERT a new tuple in a relation
- DELETE an existing tuple from a relation
- MODIFY an attribute of an existing tuple

Next slide shows an example state for the COMPANY database

# Populated database state for COMPANY

# Update Operations on Relations

---

INSERT a tuple.

DELETE a tuple.

MODIFY a tuple.

Integrity constraints should not be violated by the update operations.

Several update operations may have to be grouped together.

Updates may **propagate** to cause other updates automatically. This may be necessary to maintain integrity constraints.

# Update Operations on Relations

---

In case of integrity violation, several actions can be taken:

- Cancel the operation that causes the violation (RESTRICT or REJECT option)
- Perform the operation but inform the user of the violation
- Trigger additional updates so the violation is corrected (CASCADE option, SET NULL option)
- Execute a user-specified error-correction routine

# Possible violations for each operation

---

INSERT may violate any of the constraints:

- Domain constraint:
  - if one of the attribute values provided for the new tuple is not of the specified attribute domain
- Key constraint:
  - if the value of a key attribute in the new tuple already exists in another tuple in the relation
- Referential integrity:
  - if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation
- Entity integrity:
  - if the primary key value is null in the new tuple

# Possible violations for each operation

---

DELETE may violate only referential integrity:

- If the primary key value of the tuple being deleted is referenced from other tuples in the database
  - Can be remedied by several actions: RESTRICT, CASCADE, SET NULL (see Chapter 8 for more details)
    - RESTRICT option: reject the deletion
    - CASCADE option: propagate the new primary key value into the foreign keys of the referencing tuples
    - SET NULL option: set the foreign keys of the referencing tuples to NULL
- One of the above options must be specified during database design for each foreign key constraint

# Possible violations for each operation

---

UPDATE may violate domain constraint and NOT NULL constraint on an attribute being modified

Any of the other constraints may also be violated, depending on the attribute being updated:

- Updating the primary key (PK):
  - Similar to a DELETE followed by an INSERT
  - Need to specify similar options to DELETE
- Updating a foreign key (FK):
  - May violate referential integrity
- Updating an ordinary attribute (neither PK nor FK):
  - Can only violate domain constraints

# Summary

---

## Presented Relational Model Concepts

- Definitions
- Characteristics of relations

## Discussed Relational Model Constraints and Relational Database Schemas

- Domain constraints'
- Key constraints
- Entity integrity
- Referential integrity

## Described the Relational Update Operations and Dealing with Constraint Violations



# In-Class Exercise

(Taken from Exercise 5.15)

Consider the following relations for a database that keeps track of student enrollment in courses and the books adopted for each course:

STUDENT(SSN, Name, Major, Bdate)

COURSE(Course#, Cname, Dept)

ENROLL(SSN, Course#, Quarter, Grade)

BOOK\_ADOPTION(Course#, Quarter, Book\_ISBN)

TEXT(Book\_ISBN, Book\_Title, Publisher, Author)

**Draw a relational schema diagram specifying the foreign keys for this schema.**