

MODULE 1: APPLICATION LAYER

- 1.1 Principles of Network Applications
 - 1.1.1 Network Application Architectures
 - 1.1.1.1 Client Server Architecture
 - 1.1.1.1.1 Data Center
 - 1.1.1.2 P2P Architecture
 - 1.1.2 Processes Communicating
 - 1.1.2.1 Process
 - 1.1.2.1.1 Client & Server Processes
 - 1.1.2.1.2 Interface between the Process and the Computer Network Socket
 - 1.1.2.1.3 Addressing Processes
 - 1.1.3 Transport Services Available to Applications
 - 1.1.3.1 Reliable Data Transfer
 - 1.1.3.2 Throughput
 - 1.1.3.3 Timing
 - 1.1.3.4 Security
 - 1.1.4 Transport Services Provided by the Internet
 - 1.1.4.1 TCP Services
 - 1.1.4.2 UDP Services
- 1.2 The Web & HTTP
 - 1.2.1 Overview of HTTP
 - 1.2.1.1 Web
 - 1.2.1.2 HTTP
 - 1.2.2 Non-Persistent & Persistent Connections
 - 1.2.2.1 HTTP with Non-Persistent Connections
 - 1.2.2.2 HTTP with Persistent Connections
 - 1.2.3 HTTP Message Format
 - 1.2.3.1 HTTP Request Message
 - 1.2.3.2 HTTP Response Message
 - 1.2.4 User-Server Interaction: Cookies
 - 1.2.5 Web Caching
 - 1.2.6 The Conditional GET
- 1.3 File Transfer: FTP
 - 1.3.1 FTP Commands & Replies
- 1.4 Electronic Mail in the Internet
 - 1.4.1 SMTP
 - 1.4.2 Comparison of SMTP with HTTP
 - 1.4.3 Mail Access Protocols
 - 1.4.3.1 POP
 - 1.4.3.2 IMAP
 - 1.4.3.3 Web-Based E-Mail
- 1.5 DNS — the Internet's Directory Service
 - 1.5.1 Services Provided by DNS
 - 1.5.2 Overview of How DNS Works
 - 1.5.2.1 A Distributed, Hierarchical Database
 - 1.5.2.1.1 Recursive Queries & Iterative Queries
 - 1.5.3 DNS Records & Messages
 - 1.5.3.1 DNS Messages
- 1.6 Peer-to-Peer Applications
 - 1.6.1 P2P File Distribution
 - 1.6.1.1 BitTorrent
 - 1.6.2 Distributed Hash Table
 - 1.6.2.1 Circular Arrangement
- 1.7 Socket Programming: Creating Network Applications
 - 1.7.1 Socket Programming with UDP
 - 1.7.2 Socket Programming with TCP

MODULE 1: APPLICATION LAYER

1.1 Principles of Network Applications

- Network-applications are the driving forces for the explosive development of the internet.
- Examples of network-applications:

1) Web	5) Social networking (Facebook, Twitter)
2) File transfers	6) Video distribution (YouTube)
3) E-mail	7) Real-time video conferencing (Skype)
4) P2P file sharing	8) On-line games (World of Warcraft)
- In network-applications, program usually needs to
 - run on the different end-systems and
 - communicate with one another over the network.
- For ex: In the Web application, there are 2 different programs:
 - 1) The browser program running in the user's host (Laptop or Smartphone).
 - 2) The Web-server program running in the Web-server host.

1.1.1 Network Application Architectures

- Two approaches for developing an application:
 - 1) Client-Server architecture
 - 2) P2P (Peer to Peer) architecture

1.1.1.1 Client-Server Architecture

- In this architecture, there is a server and many clients distributed over the network (Figure 1.1a).
- The server is always-on while a client can be randomly run.
- The server is listening on the network and a client initializes the communication.
- Upon the requests from a client, the server provides certain services to the client.
- Usually, there is no communication between two clients.
- The server has a fixed IP address.
- A client contacts the server by sending a packet to the server's IP address.
- A server is able to communicate with many clients.
- The applications such as FTP, telnet, Web, e-mail etc use the client-server architecture.

1.1.1.1.1 Data Center

- Earlier, client-server architecture had a single-server host.
- But now, a single-server host is unable to keep up with all the requests from large no. of clients.
- For this reason, data-center is used.
- A data-center contains a large number of hosts.
- A data-center is used to create a powerful virtual server.
- In data center, hundreds of servers must be powered and maintained.
- For example:
 - ☐ Google has around 50 data-centers distributed around the world.
 - ☐ These 50 data-centers handle search, YouTube, Gmail, and other services.

1.1.1.2 P2P Architecture

- There is no dedicated server (Figure 1.1b).
- Pairs of hosts are called peers.
- The peers communicate directly with each other.
- The peers are not owned by the service-provider. Rather, the peers are laptops controlled by users.
- Many of today's most popular and traffic-intensive applications are based on P2P architecture.
- Examples include file sharing (BitTorrent), Internet telephone (Skype) etc.
- Main feature of P2P architectures: self-scalability.
- For ex: In a P2P file-sharing system,
 - ☐ Each peer generates workload by requesting files.
 - ☐ Each peer also adds service-capacity to the system by distributing files to other peers.
- Advantage: Cost effective. Normally, server-infrastructure & server bandwidth are not required.
- Three challenges of the P2P applications:
 - ☐ Most residential ISPs have been designed for asymmetrical bandwidth usage.
 - ☐ Asymmetrical bandwidth means there is more downstream-traffic than upstream-traffic.
 - ☐ But P2P applications shift upstream-traffic from servers to residential ISPs, which stress on the ISPs.

2) Security

- Since the highly distribution and openness, P2P applications can be a challenge to security.

3) Incentive

- Success of P2P depends on convincing users to volunteer bandwidth & resources to the applications.




Figure 1.1: (a) Client-server architecture; (b) P2P architecture

1.1.2 Processes Communicating

1.1.2.1 Process

- A process is an instance of a program running in a computer. (IPC \rightarrow inter-process communication).
- The processes may run on the 1) same system or 2) different systems.
 - 1) The processes running on the same end-system can communicate with each other using IPC.
 - 2) The processes running on the different end-systems can communicate by exchanging messages.
 - i) A sending-process creates and sends messages into the network.
 - ii) A receiving-process receives the messages and responds by sending messages back.

1.1.2.1.1 Client & Server Processes

- A network-application consists of pairs of processes:
 - 1) The process that initiates the communication is labeled as the client.
 - 2) The process that waits to be contacted to begin the session is labeled as the server.
- For example:
 - 1) In Web application, a client-browser process communicates with a Web-server-process.
 - 2) In P2P file system, a file is transferred from a process in one peer to a process in another peer.

1.1.2.1.2 Interface between the Process and the Computer Network Socket

- Any message sent from one process to another must go through the underlying-network.
- A process sends/receives message through a software-interface of underlying-network called socket.
- Socket is an API between the application-layer and the transport layer within a host (Figure 1.2).
- The application-developer has complete control at the application-layer side of the socket.
- But, the application-developer has little control of the transport-layer side of the socket.
 - 1) The choice of transport-protocol: TCP or UDP. (API \rightarrow Application Programming Interface)
 - 2) The ability to fix parameters such as maximum-buffer & maximum-segment-sizes.




Figure 1.2: Application processes, sockets, and transport-protocol

1.1.2.1.3 Addressing Processes

- To identify the receiving-process, two pieces of information need to be specified:
 - 1) IP address of the destination-host.
 - 2) Port-number that specifies the receiving-process in the destination-host.
- In the Internet, the host is identified by IP address.
- An IP address is a 32-bit that uniquely identify the host.
- Sending-process needs to identify receiving-process, a host may run several network-applications.
- For this purpose, a destination port-number is used.
- For example,
 - A Web-server is identified by port-number 80.
 - A mail-server is identified by port-number 25.

1.1.3 Transport Services Available to Applications

- Networks usually provide more than one transport-layer protocols for different applications.
- An application-developer should choose certain protocol according to the type of applications.
- Different protocols may provide different services.

1.1.3.1 Reliable Data Transfer

- Reliable means guaranteeing the data from the sender to the receiver is delivered correctly. For ex: TCP provides reliable service to an application.
- Unreliable means the data from the sender to the receiver may never arrive.
- Unreliability may be acceptable for loss-tolerant applications, such as multimedia applications.
- In multimedia applications, the lost data might result in a small glitch in the audio/video.

1.1.3.2 Throughput

- Throughput is the rate at which the sending-process can deliver bits to the receiving-process.
- Since other hosts are using the network, the throughput can fluctuate with time.
- Two types of applications:
 - 1) Bandwidth Sensitive Applications
 - These applications need a guaranteed throughput. For ex: Multimedia applications
 - Some transport-protocol provides guaranteed throughput at some specified rate (r bits/sec).
 - 2) Elastic Applications
 - These applications may not need a guaranteed throughput. For ex: Electronic mail, File transfer & Web transfers.

1.1.3.3 Timing

- A transport-layer protocol can provide timing-guarantees.
- For ex: guaranteeing every bit arrives at the receiver in less than 100 msec.
- Timing constraints are useful for real-time applications such as
 - Internet telephony

- Virtual environments
- Teleconferencing and
- Multiplayer games

1.1.3.4 Security

- A transport-protocol can provide one or more security services.
- For example,
 - 1) In the sending host, a transport-protocol can encrypt all the transmitted-data.
 - 2) In the receiving host, the transport-protocol can decrypt the received-data.

1.1.4 Transport Services Provided by the Internet

- The Internet makes two transport-protocols available to applications, UDP and TCP.
- An application-developer who creates a new network-application must use either: UDP or TCP.
- Both UDP & TCP offers a different set of services to the invoking applications.
- Table 1.1 shows the service requirements for some selected applications.

Table 1.1: Requirements of selected network-applications

Application	Data Loss	Throughput Time	Sensitive
File transfer/download	No loss	Elastic	No
E-mail	No loss	Elastic	No
Web documents	No loss	Elastic (few kbps)	No
Internet-telephony/ Video-conferencing	Loss-tolerant	Audio: few kbps–1 Mbps Video: 10 kbps–5 Mbps	Yes: 100s of ms
Streaming stored audio/video	Loss-tolerant	Same as above	Yes: few seconds
Interactive games	Loss-tolerant	Few kbps–10 kbps	Yes: 100s of ms
Instant messaging	No loss	Elastic	Yes and no

1.1.4.1 TCP Services

- An application using transport-protocol TCP, receives following 2 services.
 - ☐ Before the start of communication, client & server need to exchange control-information.
 - ☐ This phase is called handshaking phase.
 - ☐ Then, the two processes can send messages to each other over the connection.
 - ☐ After the end of communication, the applications must tear down the connection.

2) Reliable Data Transfer Service

- ☐ The communicating processes must deliver all data sent without error & in the proper order.
- TCP also includes a congestion-control.
- The congestion-control throttles a sending-process when the network is congested.

1.1.4.2 UDP Services

- UDP is a lightweight transport-protocol, providing minimal services.
- UDP is connectionless, so there is no handshaking before the 2 processes start to communicate.
- UDP provides an unreliable data transfer service.
- Unreliable means providing no guarantee that the message will reach the receiving-process.
- Furthermore, messages that do arrive at the receiving-process may arrive out-of-order.
- UDP does not include a congestion-control.
- UDP can pump data into the network-layer at any rate.

1.2 The Web & HTTP

- The appearance of Web dramatically changed the Internet.
- Web has many advantages for a lot of applications.
- It operates on demand so that the users receive what they want when they want it.
- It provides an easy way for everyone make information available over the world.
- Hyperlinks and search engines help us navigate through an ocean of Web-sites.
- Forms, JavaScript, Java applets, and many other devices enable us to interact with pages and sites.
- The Web serves as a platform for many killer applications including YouTube, Gmail, and Facebook.

1.2.1 Overview of HTTP

1.2.1.1 Web

- A web-page consists of objects (HTML → Hyper Text Markup Language).
- An object is a file such as an HTML file, a JPEG image, a Java applet, a video clip.
- The object is addressable by a single URL (URL → Uniform ResourceLocator).
- Most Web-pages consist of a base HTML file & several referenced objects.
- For example:

If a Web-page contains HTML text and five JPEG images; then the Web-page has six objects:

- 1) Base HTML file and
- 2) Five images.

- The base HTML file references the other objects in the page with the object's URLs.
- URL has 2 components:
 - 1) The hostname of the server that houses the object and
 - 2) The object's path name.
- For example:

–http://www.someSchool.edu/someDepartment/picture.gifl In above

URL,

- 1) Hostname = –www.someSchool.edu |
- 2) Path name = —/someDepartment/picture.gifl.

- The web browsers implement the client-side of HTTP. For ex: Google Chrome, Internet Explorer
- The web-servers implement the server-side of HTTP. For ex: Apache

1.2.1.2 HTTP

- HTTP is Web's application-layer protocol (Figure 1.3) (HTTP → HyperText TransferProtocol).
- HTTP defines
 - how clients request Web-pages from servers and
 - how servers transfer Web-pages to clients.

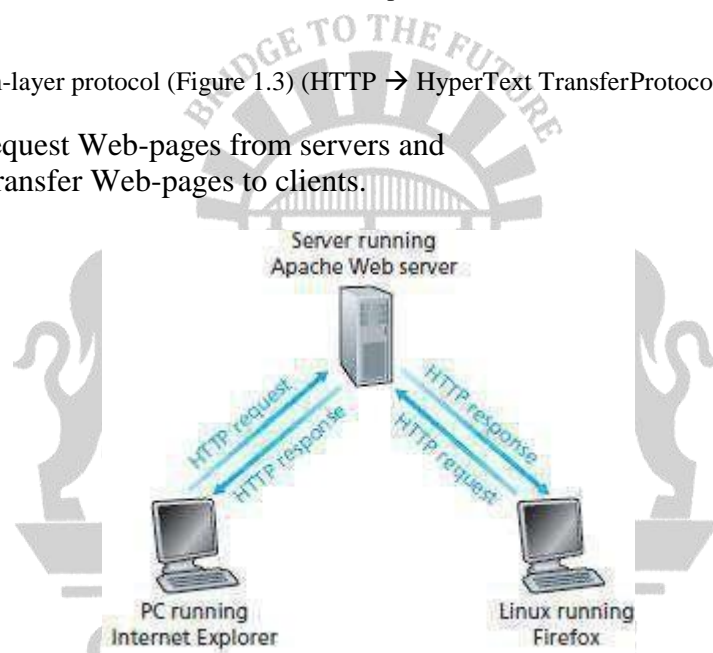


Figure 1.3: HTTP request-response behavior

- When a user requests a Web-page, the browser sends HTTP request to the server.
- Then, the server responds with HTTP response that contains the requested-objects.
- HTTP uses TCP as its underlying transport-protocol.
- The HTTP client first initiates a TCP connection with the server.
- After connection setup, the browser and the server-processes access TCP through their sockets.
- HTTP is a stateless protocol.
- Stateless means the server sends requested-object to client w/o storing state-info about the client.
- HTTP uses the client-server architecture:

- Browser that requests receive and displays Web objects.

2) Server

- Web-server sends objects in response to requests.

1.2.2 Non-Persistent & Persistent Connections

- In many internet applications, the client and server communicate for an extended period of time.
- When this client-server interaction takes place over TCP, a decision should be made:
 - 1) Should each request/response pair be sent over a separate TCP connection or
 - 2) Should all requests and their corresponding responses be sent over same TCP connection?
- These different connections are called non-persistent connections (1) or persistent connections (2).
- Default mode: HTTP uses persistent connections.

1.2.2.1 HTTP with Non-Persistent Connections

- A non-persistent connection is closed after the server sends the requested-object to the client.
- In other words, the connection is used exactly for one request and one response.
- For downloading multiple objects, multiple connections must be used.
- Suppose user enters URL:
"http://www.someSchool.edu/someDepartment/home.index"
- Assume above link contains text and references to 10 jpeg images.


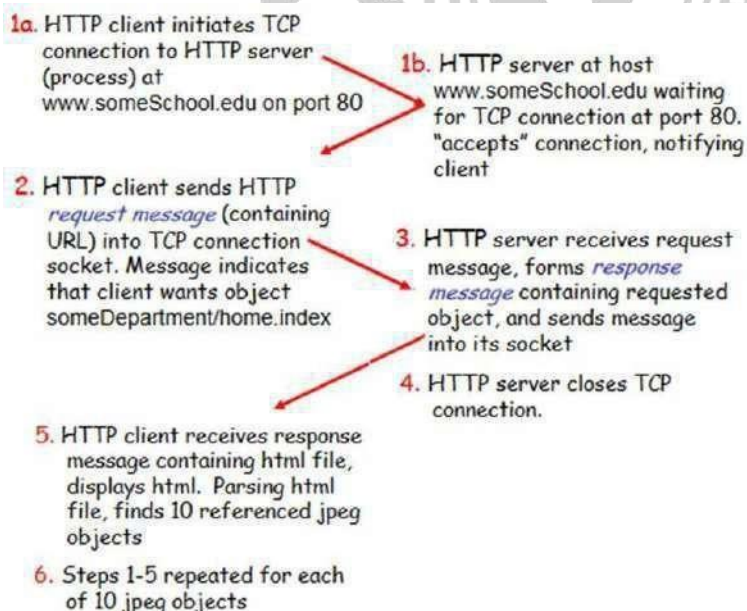


Figure 1.4: Back-of-the-envelope calculation for the time needed to request and receive an HTML file

- Here is how it works:



- RTT is the time taken for a packet to travel from client to server and then back to the client.
- The total response time is sum of following (Figure 1.4):
 - i) One RTT to initiate TCP connection (RTT \rightarrow Round TripTime).
 - ii) One RTT for HTTP request and first few bytes of HTTP response to return.
 - iii) File transmission time.
 i.e. Total response time = (i) + (ii) + (iii) = 1 RTT+ 1 RTT+ File transmission time

$$= 2(\text{RTT}) + \text{File transmission time}$$

1.2.2.2 HTTP with Persistent Connections

- Problem with Non-Persistent Connections:

- 1) A new connection must be established and maintained for each requested-object.
 - Hence, buffers must be allocated and state info must be kept in both the client and server.
 - This results in a significant burden on the server.
 - 2) Each object suffers a delivery delay of two RTTs:
 - i) One RTT to establish the TCP connection and
 - ii) One RTT to request and receive an object.
- Solution: Use persistent connections.
 - With persistent connections, the server leaves the TCP connection open after sending responses.
 - Hence, subsequent requests & responses b/w same client & server can be sent over same connection
 - The server closes the connection only when the connection is not used for a certain amount of time.
 - Default mode of HTTP: Persistent connections with pipelining.
 - Advantages:
 - 1) This method requires only one RTT for all the referenced-objects.
 - 2) The performance is improved by 20%.

1.2.3 HTTP Message Format

- Two types of HTTP messages: 1) Request-message and 2) Response-message.

1.2.3.1 HTTP Request Message

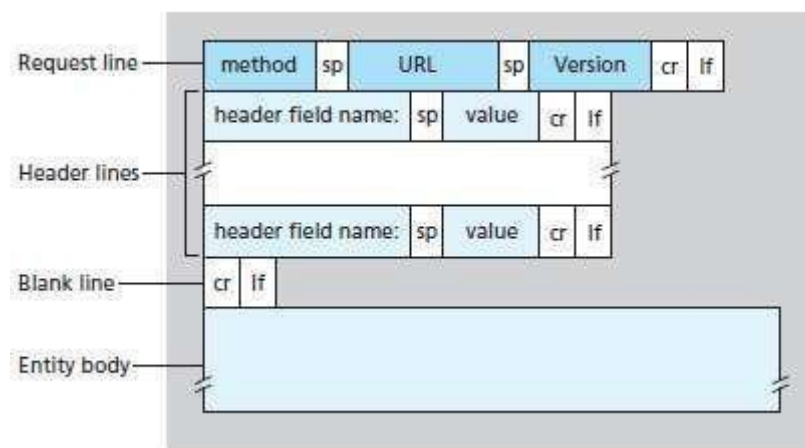


Figure 1.5: General format of an HTTP request-message

- An example of request-message is as follows: GET

```
/somedir/page.html HTTP/1.1 Host:
www.someschool.edu Connection:
close
User-agent:
Mozilla/5.0 Accept-
language: eng
```

- The request-message contains 3 sections (Figure 1.5):
 - 1) Request-line
 - 2) Header-line and
 - 3) Carriage return.
- The first line of message is called the request-line. The subsequent lines are called the header-lines.
- The request-line contains 3 fields. The meaning of the fields is as follows:
 - 1) Method
 - —GET: This method is used when the browser requests an object from the server.
 - 2) URL
 - —/somedir/page.html: This is the object requested by the browser.
 - —HTTP/1.1: This is version used by the browser.
- The request-message contains 4 header-lines. The meaning of the header-lines is as follows:
 - 1) -Host: www.someschool.edu specifies the host on which the object resides.
 - 2) -Connection: close means requesting a non-persistent connection.
 - 3) -User-agent: Mozilla/5.0 means the browser used is the Firefox.
 - 4) -Accept-language: eng means English is the preferred language.

- The method field can take following values: GET, POST, HEAD, PUT and DELETE.
 - 1) GET is used when the browser requests an object from the server.
 - 2) POST is used when the user fills out a form & sends to the server.
 - 3) HEAD is identical to GET except the server must not return a message-body in the response.
 - 4) PUT is used to upload objects to servers.
 - 5) DELETE allows an application to delete an object on a server.

1.2.3.2 HTTP Response Message

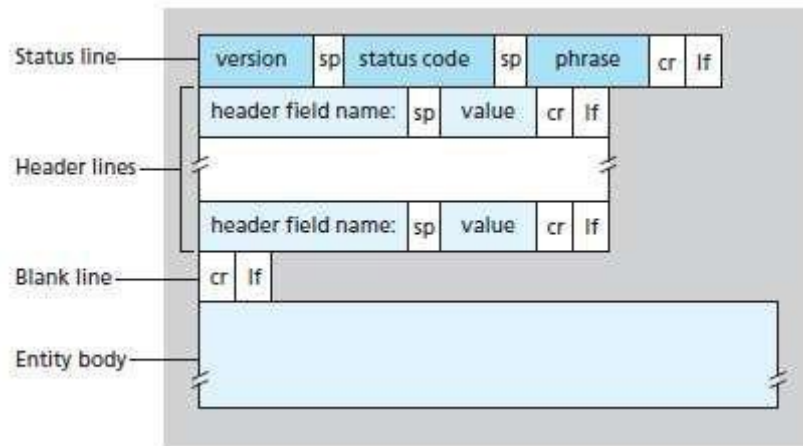


Figure 1.6: General format of an HTTP response-message

- An example of response-message is as follows:

```
HTTP/1.1 200 OK
Connection: close
Date: Tue, 09 Aug 2011 15:44:04 GMT
Server: Apache/2.2.3 (CentOS)
Last-Modified: Tue, 09 Aug 2011 15:11:03 GMT
Content-Length: 6821 Content-
Type: text/html (data data data
data data ...)
```

- The response-message contains 3 sections (Figure 1.6):
 - 1) Status line
 - 2) Header-lines and
 - 3) Data (Entity body).
- The status line contains 3 fields:
 - 1) Protocol version
 - 2) Status-code and
 - 3) Status message.
- Some common status-codes and associated messages include:
 - 1) 200 OK: Standard response for successful HTTP requests.
 - 2) 400 Bad Request: The server cannot process the request due to a client error.
 - 3) 404 Not Found: The requested resource cannot be found.
- The meaning of the Status line is as follows:
 - HTTP/1.1 200 OK!: This line indicates the server is using HTTP/1.1 & that everything is OK.
- The response-message contains 6 header-lines. The meaning of the header-lines is as follows:
 - 1) Connection: This line indicates browser requesting a non-persistent connection.
 - 2) Date: This line indicates the time & date when the response was sent by the server.
 - 3) Server: his line indicates that the message was generated by an Apache Web-server.
 - 4) Last-Modified: This line indicates the time & date when the object was last modified.
 - 5) Content-Length: This line indicates the number of bytes in the sent-object.
 - 6) Content- type: This line indicates that the object in the entity body is HTML text.

1.2.4 User-Server Interaction: Cookies

- Cookies refer to a small text file created by a Web-site that is stored in the user's computer.
- Cookies are stored either temporarily for that session only or permanently on the harddisk.
- Cookies allow Web-sites to keep track of users.

- Cookie technology has four components:
 - 1) A cookie header-line in the HTTP response-message.
 - 2) A cookie header-line in the HTTP request-message.
 - 3) A cookie file kept on the user's end-system and managed by the user's browser.
 - 4) A back-end database at the Web-site.

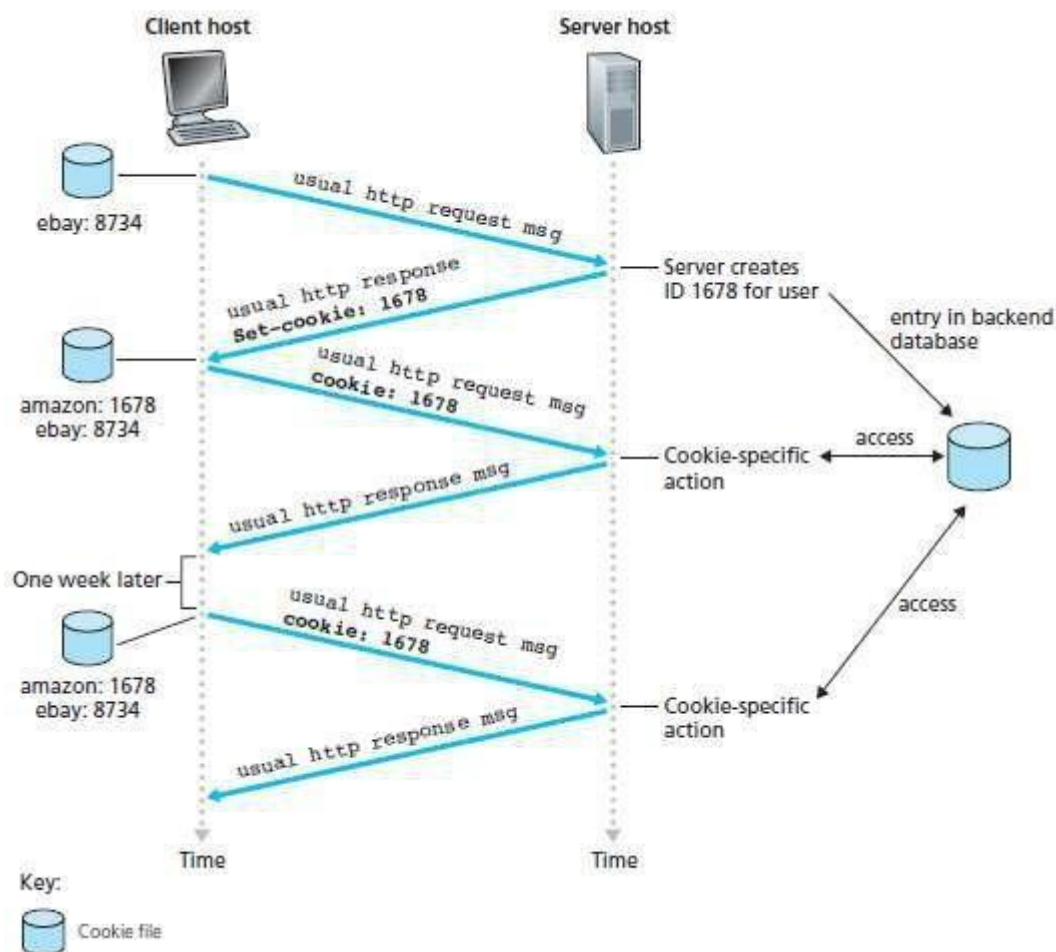


Figure 1.7: Keeping user state with cookies

- Here is how it works (Figure 1.7):
 - 1) When a user first time visits a site, the server
 - creates a unique identification number (1678) and
 - creates an entry in its back-end database by the identification number.
 - 2) The server then responds to user's browser.
 - HTTP response includes Set-cookie: header which contains the identification number (1678)
 - 3) The browser then stores the identification number into the cookie-file.
 - 4) Each time the user requests a Web-page, the browser
 - extracts the identification number from the cookie file, and
 - puts the identification number in the HTTP request.
 - 5) In this manner, the server is able to track user's activity at the web-site.

1.2.5 Web Caching

- A Web-cache is a network entity that satisfies HTTP requests on the behalf of an original Web-server.
- The Web-cache has disk-storage.
- The disk-storage contains copies of recently requested-objects.




Figure 1.8: Clients requesting objects through a Web-cache (or Proxy server)

- Here is how it works (Figure 1.8):
 - 1) The user's HTTP requests are first directed to the web-cache.
 - 2) If the cache has the object requested, the cache returns the requested-object to the client.
 - 3) If the cache does not have the requested-object, then the cache
 - connects to the original server and
 - asks for the object.
 - 4) When the cache receives the object, the cache
 - stores a copy of the object in local-storage and
 - sends a copy of the object to the client.
- A cache acts as both a server and a client at the same time.
 - 1) The cache acts as a server when the cache
 - receives requests from a browser and
 - sends responses to the browser.
 - 2) The cache acts as a client when the cache
 - requests to an original server and
 - receives responses from the origin server.
- Advantages of caching:
 - 1) To reduce response-time for client-request.
 - 2) To reduce traffic on an institution's access-link to the Internet.
 - 3) To reduce Web-traffic in the Internet.

1.2.6 The Conditional GET

- Conditional GET refers a mechanism that allows a cache to verify that the objects are up to date.
- An HTTP request-message is called conditional GET if
 - 1) Request-message uses the GET method and
 - 2) Request-message includes an If-Modified-Since: header-line.
- The following is an example of using conditional GET: GET

```
/fruit/kiwi.fig HTTP/1.1
Host: www.exoriguecuisine.com
If-modified-since: Wed, 7 Sep 2011 09:23:24
```

- The response is:

```
HTTP/1.1 304 Not Modified
Date: Sat, 15 Oct 2011 15:39:29
```

1.3 File Transfer: FTP

- FTP is used by the local host to transfer files to or from a remote-host over the network.
- FTP uses client-server architecture (Figure 1.9).
- FTP uses 2 parallel TCP connections (Figure 1.10):
 - 1) Control Connection
 - ☐ The control-connection is used for sending control-information b/w local and remote-hosts.
 - ☐ The control-information includes:
 - user identification
 - password

- commands to change directory and
- commands to put & get files.

2) Data Connection

- The data-connection is used to transfer files.




Figure 1.9: FTP moves files between local and remote file systems



Figure 1.10: Control and data-connections

- Here is how it works:
 - 1) When session starts, the client initiates a control-connection with the server on port 21.
 - 2) The client sends user-identity and password over the control-connection.
 - 3) Then, the server initiates data-connection to the client on port 20.
 - 4) FTP sends exactly one file over the data-connection and then closes the data-connection.
 - 5) Usually, the control-connection remains open throughout the duration of the user-session.
 - 6) But, a new data-connection is created for each file transferred within a session.
- During a session, the server must maintain the state-information about the user.
- For example:
 - The server must keep track of the user's current directory.
- Disadvantage:
 - Keeping track of state-info limits the no. of sessions maintained simultaneously by a server.

1.3.1 FTP Commands & Replies

- The commands are sent from client to server.
- The replies are sent from server to client.
- The commands and replies are sent across the control-connection in 7-bit ASCII format.
- Each command consists of 4-uppercase ASCII characters followed by optional arguments.
- For example:
 - 1) **USER** username
 - Used to send the user identification to the server.
 - 2) **PASS** password
 - Used to send the user password to the server.
 - 3) **LIST**
 - Used to ask the server to send back a list of all the files in the currentremote directory.
 - 4) **RETR** filename
 - Used to retrieve a file from the current directory of the remote-host.
 - 5) **STOR** filename
 - Used to store a file into the current directory of the remote-host.
- Each reply consists of 3-digit numbers followed by optional message.
- For example:
 - 1) 331 Username OK, password required
 - 2) 125 Data-connection already open; transfer starting
 - 3) 425 Can't open data-connection
 - 4) 452 Error writing file

1.4 Electronic Mail in the Internet

- e-mail is an asynchronous communication medium in which people send and read messages.
- e-mail is fast, easy to distribute, and inexpensive.
- e-mail has features such as
 - messages with attachments
 - hyperlinks
 - HTML-formatted text and
 - embedded photos.
- Three major components of an e-mail system (Figure 1.11):

- User-agents allow users to read, reply to, forward, save and compose messages.
- For example: Microsoft Outlook and Apple Mail
- Mail-servers contain mailboxes for users.
- A message is first sent to the sender's mail-server.
- Then, the sender's mail-server sends the message to the receiver's mail-server.
- If the sender's server cannot deliver mail to receiver's server, the sender's server
 - holds the message in a message queue and
 - attempts to transfer the message later.

3) SMTP (Simple Mail Transfer Protocol)

- SMTP is an application-layer protocol used for email.
- SMTP uses TCP to transfer mail from the sender's mail-server to the recipient's mail-server.
- SMTP has two sides:
 - 1) A client-side, which executes on the sender's mail-server.
 - 2) A server-side, which executes on the recipient's mail-server.
- Both the client and server-sides of SMTP run on every mail-server.
- When a mail-server receives mail from other mail-servers, the mail-server acts as a server. When a mail-server sends mail to other mail-servers, the mail-server acts as a client.

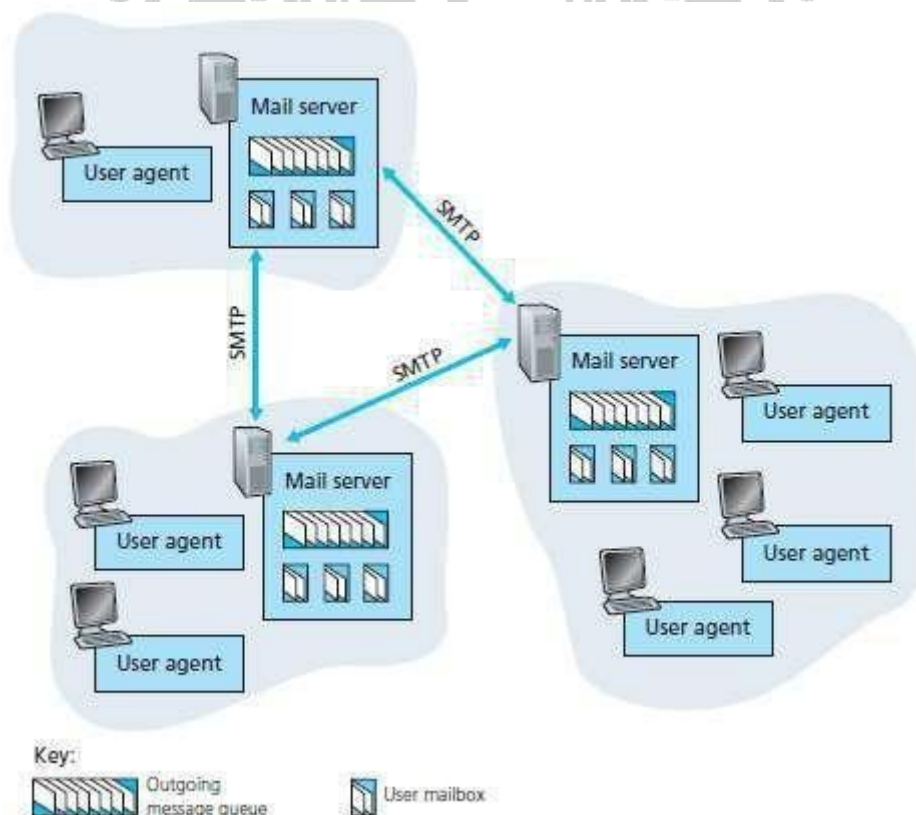


Figure 1.11: A high-level view of the Internet e-mail system

1.4.1 SMTP

- SMTP is the most important protocol of the email system.
- Three characteristics of SMTP (that differs from other applications):

- 1) Message body uses 7-bit ASCII code only.
 - 2) Normally, no intermediate mail-servers used for sending mail.
 - 3) Mail transmissions across multiple networks through mail relaying.
- Here is how it works:
 - 1) Usually, mail-servers are listening at port 25.
 - 2) The sending server initiates a TCP connection to the receiving mail-server.
 - 3) If the receiver's server is down, the sending server will try later.
 - 4) If connection is established, the client & the server perform application-layer handshaking.
 - 5) Then, the client indicates the e-mail address of the sender and the recipient.
 - 6) Finally, the client sends the message to the server over the same TCP connection.

1.4.2 Comparison of SMTP with HTTP

- 1) HTTP is mainly a pull protocol. This is because
 - someone loads information on a web-server and
 - users use HTTP to pull the information from the server.
- On the other hand, SMTP is primarily a push protocol. This is because
 - the sending mail-server pushes the file to receiving mail-server.
- 2) SMTP requires each message to be in seven-bit ASCII format.
 - If message contains binary-data, the message has to be encoded into 7-bit ASCII format.
 - HTTP does not have this restriction.
- 3) HTTP encapsulates each object of message in its own response-message.
 - SMTP places all of the message's objects into one message.

1.4.3 Mail Access Protocols

- It is not realistic to run the mail-servers on PC & laptop. This is because
 - mail-servers must be always-on and
 - mail-servers must have fixed IP addresses
- Problem: How a person can access the email using PC or laptop?
- Solution: Use mail access protocols.
- Three mail access protocols:
 - 1) Post Office Protocol (POP)
 - 2) Internet Mail Access Protocol (IMAP) and
 - 3) HTTP.

1.4.3.1 POP

- POP is an extremely simple mail access protocol.
- POP server will listen at port 110.
- Here is how it works:
 - ☐ The user-agent at client's computer opens a TCP connection to the main server.
 - ☐ POP then progresses through three phases:
 - ☐ The user-agent sends a user name and password to authenticate the user.
 - 2) Transaction
 - ☐ The user-agent retrieves messages.
 - ☐ Also, the user-agent can
 - mark messages for deletion
 - remove deletion marks &
 - obtain mail statistics.
 - ☐ The user-agent issues commands, and the server responds to each command with a reply.
 - ☐ There are two responses:
 - i) +OK: used by the server to indicate that the previous command was fine.
 - ii) -ERR: used by the server to indicate that something is wrong.
 - ☐ After user issues a quit command, the mail-server removes all messages marked for deletion.
- Disadvantage:

The user cannot manage the mails at remote mail-server. For ex: user cannot delete messages.

1.4.3.2 IMAP

- IMAP is another mail access protocol, which has more features than POP.
- An IMAP server will associate each message with a folder.

- When a message first arrives at server, the message is associated with recipient's INBOX folder
- Then, the recipient can
 - move the message into a new, user-created folder
 - read the message
 - delete the message and
 - search remote folders for messages matching specific criteria.
- An IMAP server maintains user state-information across IMAP sessions.
- IMAP permits a user-agent to obtain components of messages.
For example, a user-agent can obtain just the message header of a message.

1.4.3.3 Web-Based E-Mail

- HTTPs are now used for Web-based email accessing.
- The user-agent is an ordinary Web browser.
- The user communicates with its remote-server via HTTP.
- Now, Web-based emails are provided by many companies including Google, Yahoo etc.

1.5 DNS — The Internet's Directory Service

- DNS is an internet service that translates domain-names into IP addresses.
For ex: the domain-name —www.google.com might translate to IP address —198.105.232.41.
- Because domain-names are alphabetic, they are easier to remember for human being.
- But, the Internet is really based on IP addresses (DNS → Domain Name System).

1.5.1 Services Provided by DNS

- The DNS is
 - 1) A distributed database implemented in a hierarchy of DNS servers.
 - 2) An application-layer protocol that allows hosts to query the distributed database.
- DNS servers are often UNIX machines running the BIND software.
- The DNS protocol runs over UDP and uses port 53. (BIND → Berkeley Internet Name Domain)
- DNS is used by application-layer protocols such as HTTP, SMTP, and FTP.
- Assume a browser requests the URL www.someschool.edu/index.html.
- Next, the user's host must first obtain the IP address of www.someschool.edu
- This is done as follows:
 - 1) The same user machine runs the client-side of the DNS application.
 - 2) The browser
 - extracts the hostname —www.someschool.edu from the URL and
 - passes the hostname to the client-side of the DNS application.
 - 3) The client sends a query containing the hostname to a DNS server.
 - 4) The client eventually receives a reply, which includes the IP address for the hostname.
 - 5) After receiving the IP address, the browser can initiate a TCP connection to the HTTP server.
- DNS also provides following services:
 - 1) Host Aliasing
 - A host with a complicated hostname can have one or more alias names.
 - 2) Mail Server Aliasing
 - For obvious reasons, it is highly desirable that e-mail addresses be mnemonic.
 - 3) Load Distribution
 - DNS is also used to perform load distribution among replicated servers.
 - Busy sites are replicated over multiple servers & each server runs on a different system.

1.5.2 Overview of How DNS Works

- Distributed database design is more preferred over centralized design because:
 - 1) A Single Point of Failure
 - If the DNS server crashes then the entire Internet will not stop.
 - 2) Traffic Volume
 - A Single DNS Server cannot handle the huge global DNS traffic.
 - But with distributed system, the traffic is distributed and reduces overload on server.
 - 3) Distant Centralized Database
 - A single DNS server cannot be —close to all the querying clients.
 - If we put the single DNS server in Mysore, then all queries from USA must travel to the other side of the globe.
 - This can lead to significant delays.

4) Maintenance

- The single DNS server would have to keep records for all Internet hosts.
- This centralized database has to be updated frequently to account for every new host.

1.5.2.1 A Distributed, Hierarchical Database

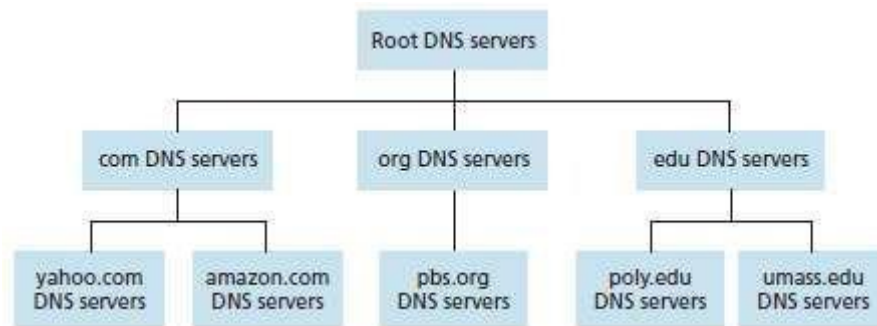


Figure 1.12: Portion of the hierarchy of DNS servers

- Suppose a client wants to determine IP address for hostname —www.amazon.com (Figure 1.12):
 - 1) The client first contacts one of the root servers, which returns IP addresses for TLD servers
 - 2) Then, the client contacts one of these TLD servers.
 - The TLD server returns the IP address of an authoritative-server for —amazon.com.
 - 3) Finally, the client contacts one of the authoritative-servers for amazon.com.
 - The authoritative-server returns the IP address for the hostname —www.amazon.com.

1.5.2.1.1 Recursive Queries & Iterative Queries




Figure 1.13: Interaction of the various DNS servers

- The example shown in Figure 1.13 makes use of both recursive queries and iterative queries.
- The query 1 sent from cis.poly.edu to dns.poly.edu is a recursive query. This is because
 - the query asks dns.poly.edu to obtain the mapping on its behalf.
- But the subsequent three queries 2, 4 and 6 are iterative. This is because
 - all replies are directly returned to dns.poly.edu.

1.5.3 DNS Records & Messages

- The DNS server stores resource-records (RRs).
- RRs provide hostname-to-IP address mappings.
- Each DNS reply message carries one or more resource-records.
- A resource-record is a 4-tuple that contains the following fields: (Name, Value, Type, TTL)
- TTL (time to live) determines when a resource should be removed from a cache.
- The meaning of Name and Value depend on Type:
 - 1) If Type=A, then Name is a hostname and Value is the IP address for the hostname.
 - Thus, a Type A record provides the standard hostname-to-IP address mapping. For ex: (relay1.bar.foo.com, 145.37.93.126, A)
 - 2) If Type=NS, then
 - i) Name is a domain (such as foo.com) and
 - ii) Value is the hostname of an authoritative DNS server.
 - This record is used to route DNS queries further along in the query chain.
 - 3) If Type=CNAME, then Value is a canonical hostname for the alias hostname Name.
 - This record can provide querying hosts the canonical name for a hostname. For ex: (foo.com, relay1.bar.foo.com, CNAME) is a CNAME record.
 - 4) If Type=MX, Value is the canonical name of a mail-server that has an alias hostname Name. ➤
 - MX records allow the hostnames of mail-servers to have simple aliases.
 - For ex: (foo.com, mail.bar.foo.com, MX) is an MX record.

1.5.3.1 DNS Messages

- Two types of DNS messages: 1) query and 2) reply.
- Both query and reply messages have the same format.

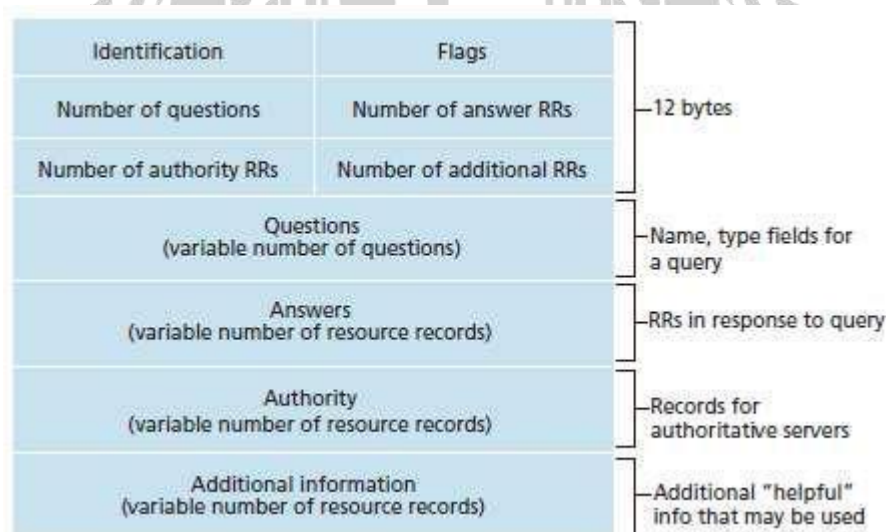


Figure 1.14: DNS message format

- The various fields in a DNS message are as follows (Figure 1.14):

1) Header Section

- The first 12 bytes is the header-section.
- This section has following fields:
 - i) Identification
 - This field identifies the query.
 - This identifier is copied into the reply message to a query.
 - This identifier allows the client to match received replies with sent queries.
 - ii) Flag
 - This field has following 3 flag-bits:

a) Query/Reply

- ✕ This flag-bit indicates whether the message is a query (0) or a reply (1).

b) Authoritative

✕ This flag-bit is set in a reply message when a DNS server is an authoritative-server.

c) Recursion Desired

✕ This flag-bit is set when a client desires that the DNS server perform recursion.

iii) Four Number-of-Fields

□ These fields indicate the no. of occurrences of 4 types of data sections that follow the header.

2) Question Section

- This section contains information about the query that is being made.
- This section has following fields:

i) Name

□ This field contains the domain-name that is being queried.

ii) type

□ This field indicates the type of question being asked about the domain-name.

3) Answer Section

- This section contains a reply from a DNS server.
- This section contains the resource-records for the name that was originally queried.
- A reply can return multiple RRs in the answer, since a hostname can have multiple IP addresses.

4) Authority Section

- This section contains records of other authoritative-servers.
- This section contains other helpful records.

1.6 Peer-to-Peer Applications

- Peer-to-peer architecture is different from client-server architecture.
- In P2P, each node (called peers) acts as a client and server at the same time.
- The peers are not owned by a service-provider.
- The peers are not supposed to be always listening on the Internet.
- The peers are dynamic, i.e., some peers will join some peers will leave from time to time.

1.6.1 P2P File Distribution

- One popular P2P file distribution protocol is BitTorrent
- Consider the following scenarios:

Suppose a server has a large file and N computers want to download the file (Figure 1.15).

- 1) In client-server architecture, each of the N computers will
 - connect to the server &
 - download a copy of the file to local-host.
- 2) In P2P architecture, a peer need not necessarily download a copy from the server. ➤
Rather, the peer may download from other peers.




Figure 1.15: An illustrative file distribution problem

- Let us define the following: Length of the file = F
- Upload rate for the server = u_s

rate for i th computer = u_i Download-rate
for i th computer = d_i

Distribution-time for the client-server architecture = D_{cs}

Server needs transmit = N_F bits.

Lowest download-rate of peer = d_{\min}

Distribution-time for the P2P architecture =

D_{P2P}

Case 1: Client-Server Architecture

- We have 2 observations:
 - The server must transmit one copy of the file to each of the N peers.
 - Since the server's upload rate is u_s , the distribution-time is at least N_F/u_s .
 - The peer with the lowest download-rate cannot obtain all F bits of the file in less than F/d_{\min} .
 - Thus, the minimum distribution-time is at least F/d_{\min} .
- Putting above 2 observations together, we have

$$D_{cs} = \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{\min}} \right\}$$

Case 2: P2P Architecture

- We have 3 observations:
 - At the beginning of the distribution, only the server has the file.
 - So, the minimum distribution-time is at least F/u_s .
 - The peer with the lowest download-rate cannot obtain all F bits of the file in less than F/d_{\min} .
 - Thus, the minimum distribution-time is at least F/d_{\min} .
 - The total upload capacity of the system as a whole is $u_{\text{total}} = u_s + u_1 + u_2 + \dots + u_N$.
 - The system must deliver F bits to each of the N peers.
 - Thus, the minimum distribution-time is at least $NF/(u_s + u_1 + u_2 + \dots + u_N)$.
- Putting above 3 observations together, we have

$$D_{P2P} \geq \max \left\{ \frac{F}{u_s}, \frac{F}{d_{\min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

- Figure 1.16 compares the minimum distribution-time for the client-server and P2P architectures.

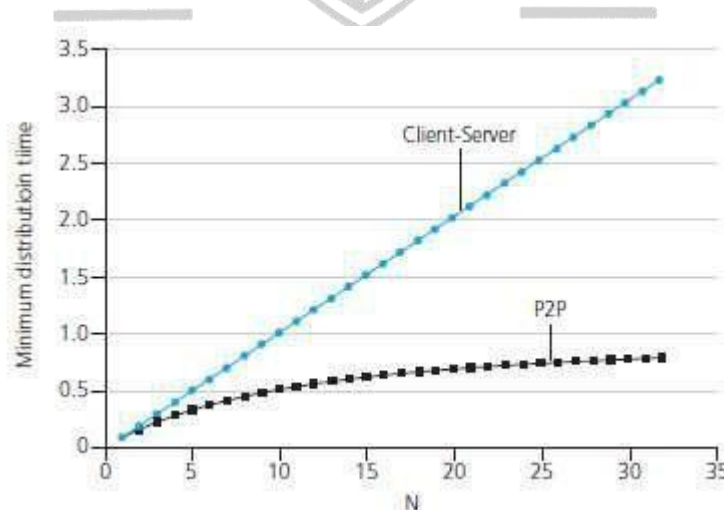


Figure 1.16: Distribution-time for P2P and client-server architectures

- The above comparison shows that when N is large,
 - P2P architecture is consuming less distribution-time.
 - P2P architecture is self-scaling.

1.6.1.1 BitTorrent

- The collection of all peers participating in the distribution of a particular file is called a torrent.
- Peers download equal-size chunks of the file from one another. Chunk size = 256 KBytes.

- The peer also uploads chunks to other peers.
- Once a peer has acquired the entire file, the peer may leave the torrent or remain in the torrent.
- Each torrent has an infrastructure node called tracker.
- Here is how it works (Figure 1.17):
 - 1) When a peer joins a torrent, the peer
 - registers itself with the tracker and
 - periodically informs the tracker that it is in the torrent.
 - 2) When a new peer joins the torrent, the tracker
 - randomly selects a subset of peers from the set of participating peers and
 - sends the IP addresses of these peers to the new peer.
 - 3) Then, the new peer tries to establish concurrent TCP connections with all peers on this list. ➤ All peers on the list are called neighboring-peers.
 - 4) Periodically, the new peer will ask each of the neighboring-peers for the set of chunks.
- To choose the chunks to download, the peer uses a technique called rarest-first.
- Main idea of rarest-first:
 - Determine the chunks that are the rarest among the neighbors and
 - Request then those rarest chunks first.




Figure 1.17: File distribution with BitTorrent

(Source : DIGINOTES)

1.6.2 Distributed Hash Table

- We can use P2P architecture to form a distributed database.
- We consider a simple database, which contains (key, value) pairs.
- Each peer will only hold a small subset of the data.
- Any peer can query the distributed database with a particular key.
- Then, the database will
 - locate the peers that have the corresponding (key, value) pairs and
 - return the key-value to the querying peer.
- Any peer will also be allowed to insert new key-value pairs into the database.
- Such a distributed database is referred to as a distributed hash table (DHT).
- To construct the database, we need to use some hash-functions.
- The input of a hash-function can be a large number.
- The outline of building a DHT is as follows:
 - 1) Assign an identifier to each peer, where the identifier is an n-bit string.
 - So, we can view the identifier as an integer at the range from 0 to $2^n - 1$.
 - 2) For a data pair (key, value), the hash value of the key is computed.
 - Then the data is stored in the peer whose identifier is closest to the key.
 - 3) To insert or retrieve data, first we need to find the appropriate peer.

- Problem: It is not realistic to let the peer to store all of the other peer's identifiers. Solution: Use a circular arrangement.




Figure 1.18: (a) A circular DHT. Peer 3 wants to determine who is responsible for key 11. (b) A circular DHT with shortcuts

1.6.2.1 Circular Arrangement

- In this example, the identifiers are range [0, 15] (4-bit strings) and there are eight peers.
- Each peer is only aware of its immediate successor and predecessor (Figure 1.18a).
- So, each peer just needs to track two neighbors.
- When a peer asks for a key, the peer sends the message clockwise around the circle.
- For example:
 - The peer 4 sends a message saying "Who is responsible for key 11?"
 - he message will forward through peers 5, 8, 10 and reach peer 12.
 - he peer 12 determines that it is the closest peer to key 11.
 - At this point, peer 12 sends a message back to the querying peer 4.
- But when the circle is too large, a message may go through a large number of peers to get answer.
- Problem: There is trade off between
 - i) Number of neighbors each peer has to track and
 - ii) Number of message to be sent to resolve a single query.
- Solution: To reduce the query time, add "shortcuts" to the circular arrangement (Figure 1.18b).

1.7 Socket Programming: Creating Network Applications

- Two types of network-applications:
 - 1) First type is an implementation whose operation is specified in a protocol standard (RFC)
 - Such an application is referred to as "openl".
 - The client & server programs must conform to the rules dictated by the RFC.
 - 2) Second type is a proprietary network-application.
 - The client & server programs use an application-layer protocol not openly published in a RFC.
 - A single developer creates both the client and server programs.
 - The developer has complete control over the code.
- During development phase, the developer must decide whether the application uses TCP or UDP.

1.7.1 Socket Programming with UDP

- Consider client-server application in which following events occurs:
 - 1) The client
 - reads a line of characters (data) from the keyboard and
 - sends the data to the server.
 - 2) The server receives the data and converts the characters to uppercase.
 - 3) The server sends the modified data to the client.
 - 4) The client receives the modified data and displays the line on its screen.

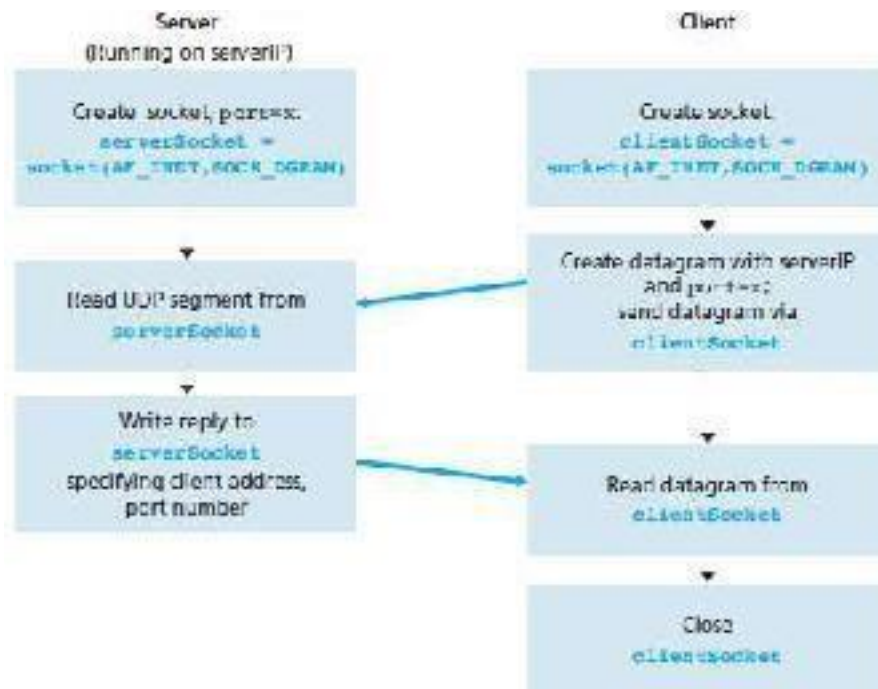


Figure 1.19: The client-server application using UDP

- The client-side of the application is as follows (Figure 1.19): from socket

```

import * //This line declares socket within the program. serverName =
„hostname“ // This line sets the server name to “hostname”.
serverPort = 12000 // This line sets the server port# to “12000”.
clientSocket = socket(socket.AF_INET, socket.SOCK_DGRAM) //This line creates the
client’s socket
message = raw_input(" Input lowercase sentence: ") //This line reads
the data at the client
clientSocket.sendto(message, (serverName, serverPort)) //This
line sends data into the socket
modifiedMessage, serverAddress =
  
```

Here, AF_INET indicates address family

SOCK_DGRAM indicates UDP as socket type contains server’s IP address & port#

- The server-side of the application is as

```

follows: from socket
import *
serverPort = 12000
serverSocket = socket(AF_INET, SOCK_DGRAM)
serverSocket.bind(("", serverPort)) // This line assigns the port# 12000 to the
server’s socket.
print "The server is ready to receive"
while 1:
    message, clientAddress = serverSocket.recvfrom(2048)
  
```

1.7.2 Socket Programming with TCP

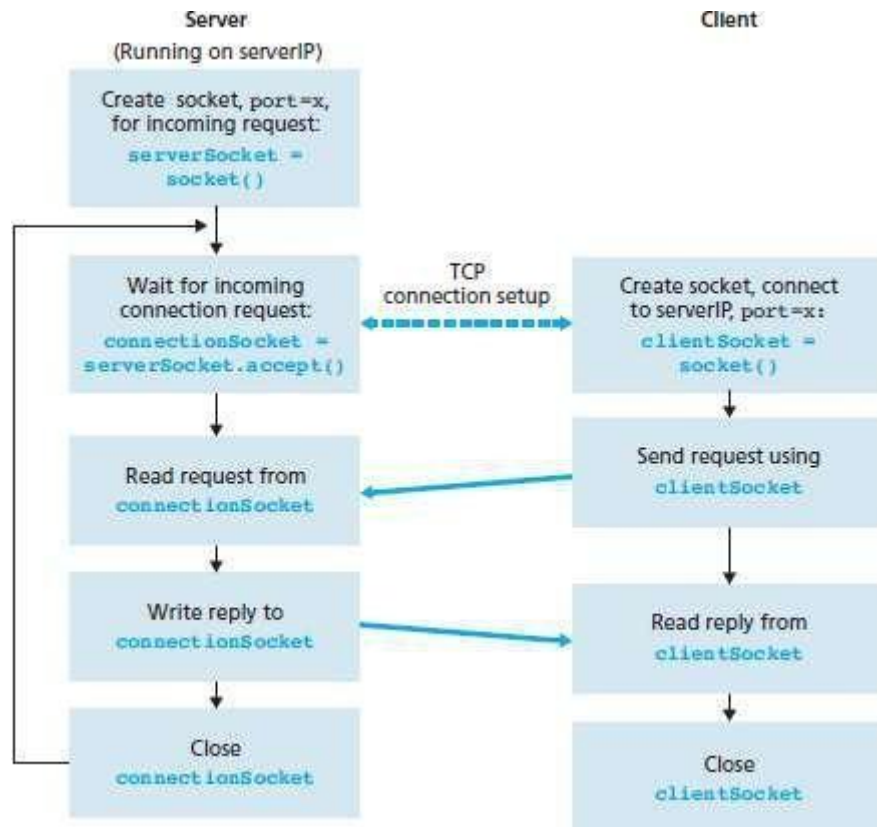


Figure 1.20: The client-server application using TCP

- The client-side of the application is as follows (Figure 1.20):

```

from socket import *
serverName =
"servername"
serverPort = 12000
clientSocket = socket(AF_INET, SOCK_STREAM)
clientSocket.connect((serverName,serverPort)) // This line initiates TCP connection
b/w client & server
sentence = raw_input("Input lowercase sentence:")
clientSocket.send(sentence)
modifiedSentence =
  
```

- The server-side of the application is as follows:

```

from socket
import *
serverPort =
12000
serverSocket = socket(AF_INET,SOCK
STREAM)
serverSocket.bind(("",serverPort))
serverSocket.listen(1) // This line specifies no. of connection-requests from the client to
server
print "The server is ready to
  
```